Big Data Streaming Pipeline and Integration Platform Project

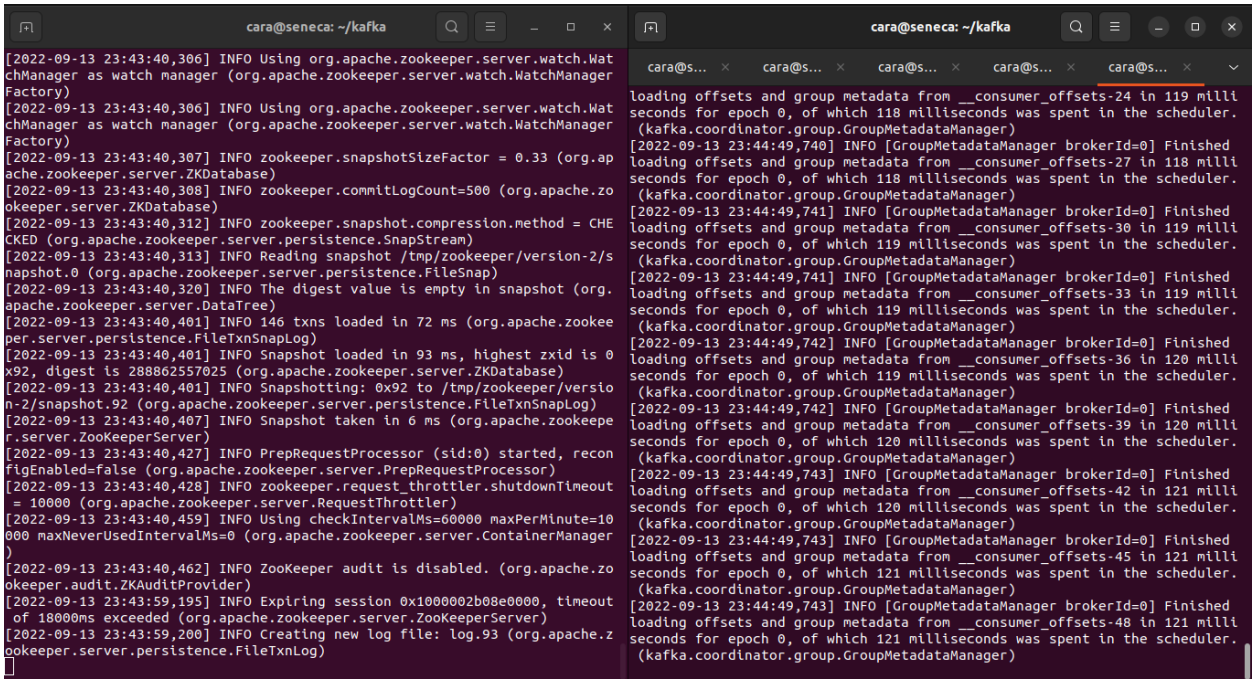

Zookeeper and Kakfa Running

```
## run zookeeper
bin/zookeeper-server-start.sh config/zookeeper.properties

##run kafka
JMX_PORT=8004 bin/kafka-server-start.sh config/server.properties
```

Commands used to run zookeeper and kafka.

Part 1: 10 snippets of producer's code

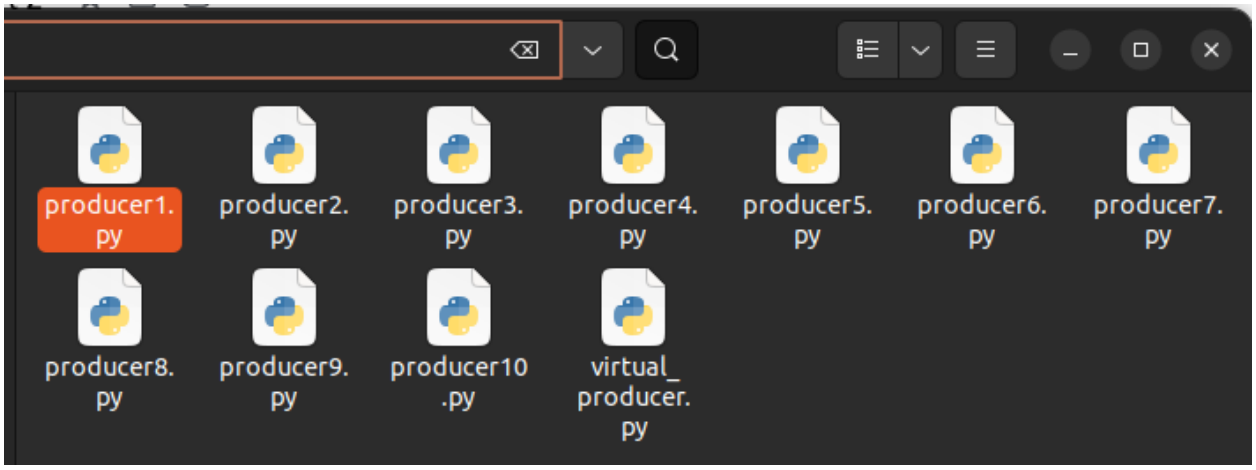

10 instances of producers, each one has unique sensor_id from 1 -10

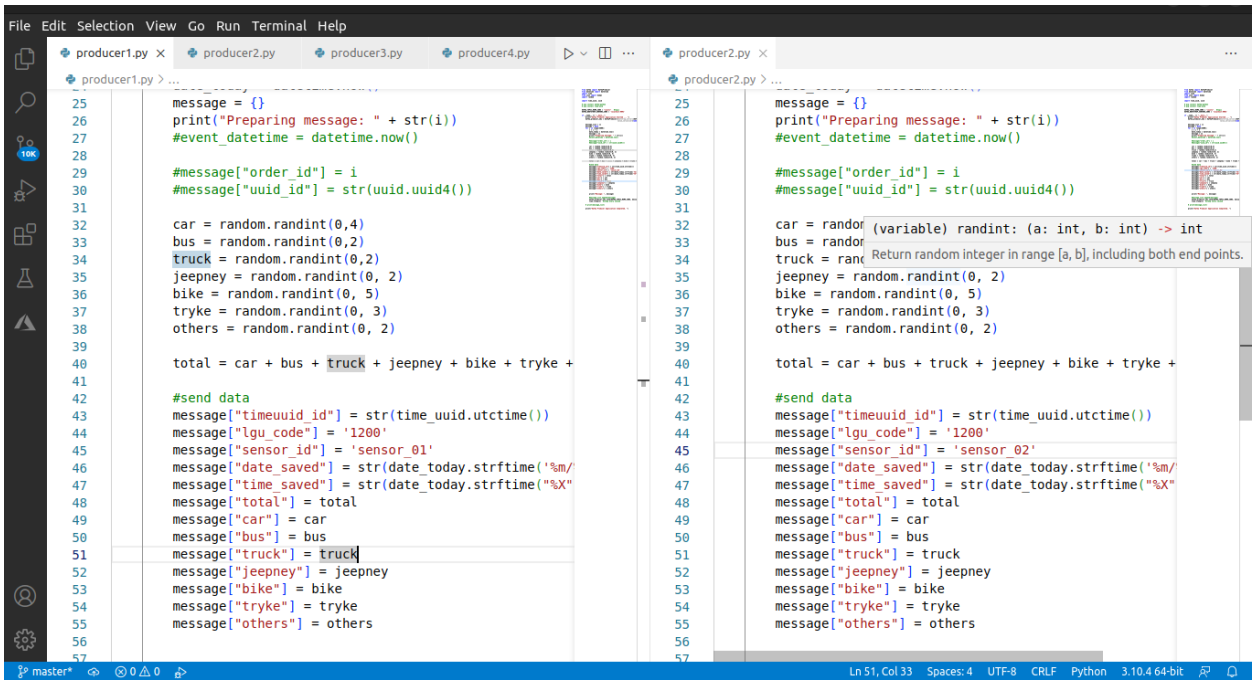

Producer 1 and 2

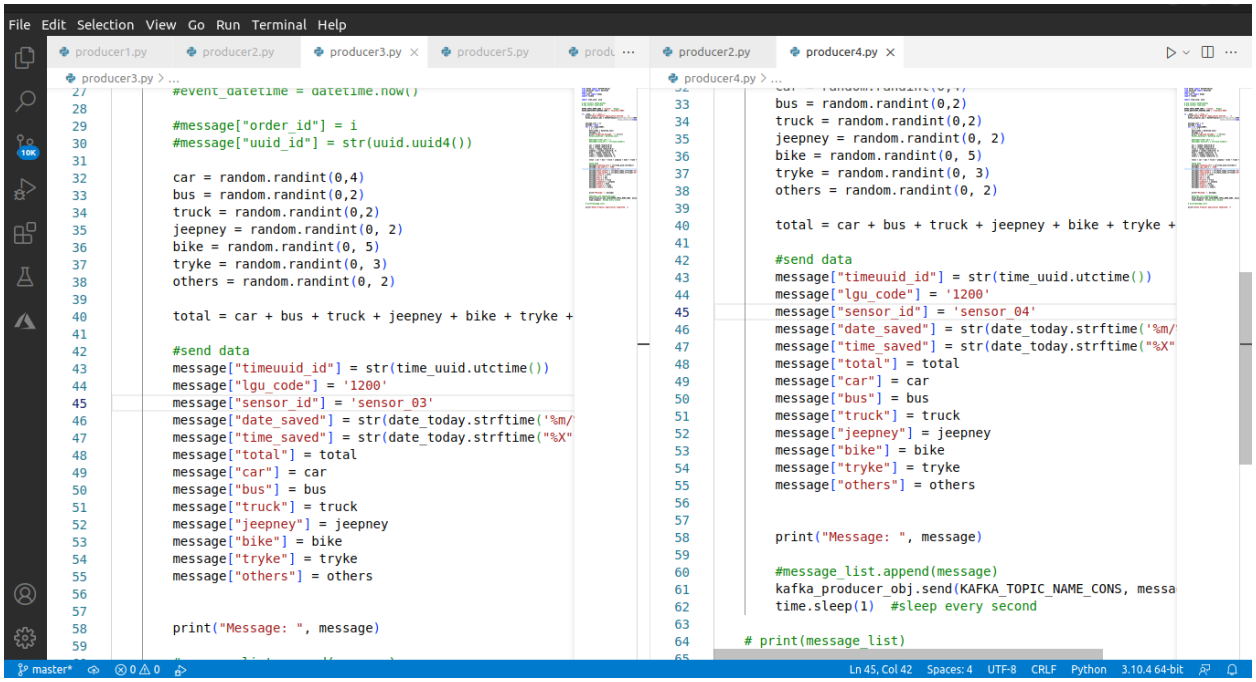

Producer 3 and 4

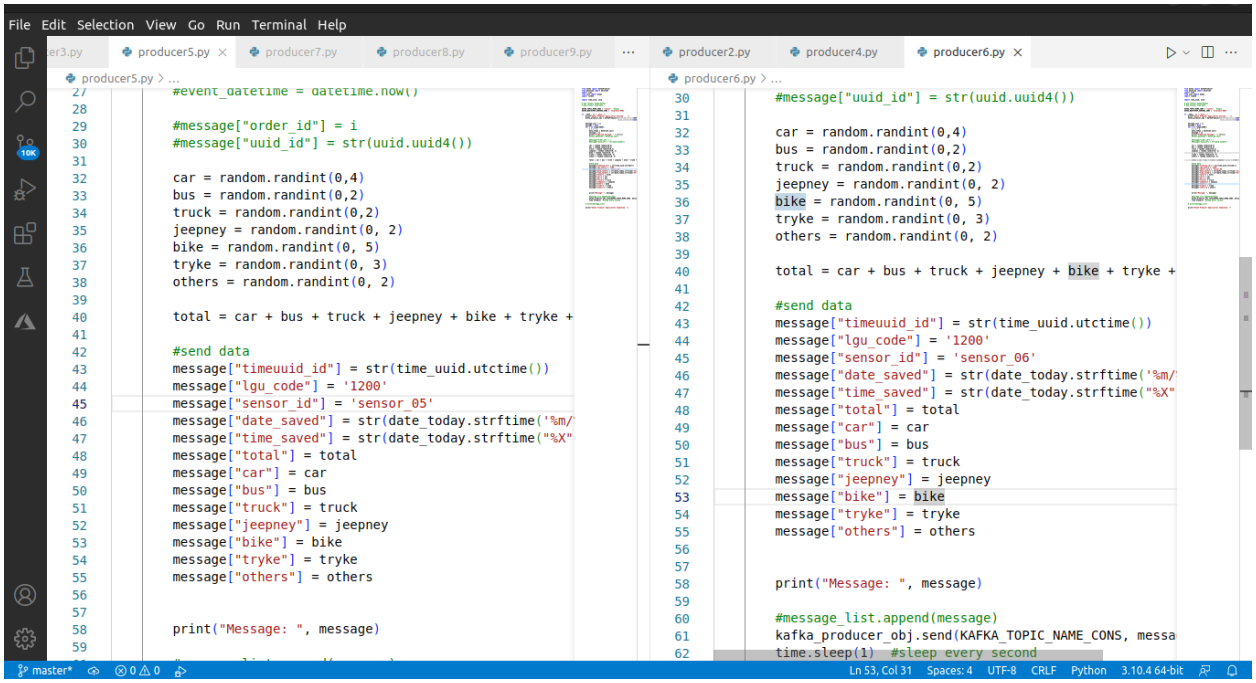

Producer 5 and 6

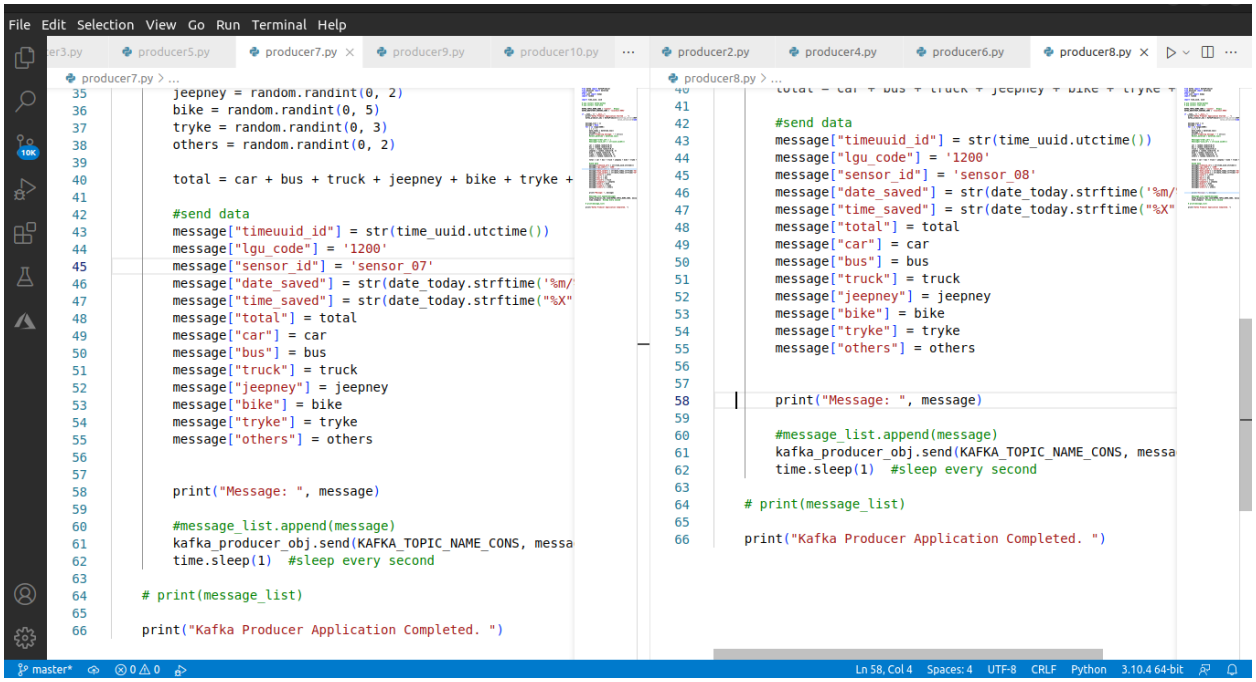

Producer 7 and 8

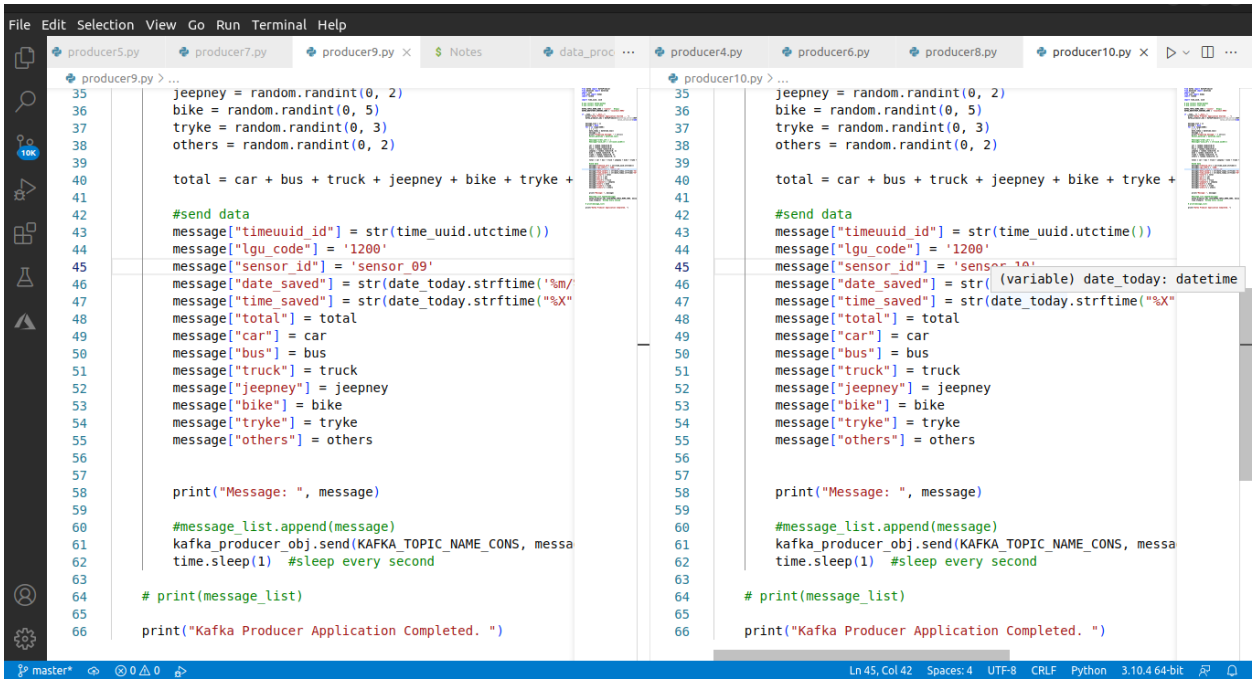

Producer 9 and 10

Part II. Create topic and run producers

```
Topic: counts   TopicId: ShruQO1XSDSWG6WM_wR0fw PartitionCount: 1       Replicat
ionFactor: 1    Configs: segment.bytes=1073741824
        Topic: counts   Partition: 0    Leader: 0       Replicas: 0     Isr: 0
```

Created topic ‘counts’. Different producers

```
KAFKA_TOPIC_NAME_CONS = "counts"    #topic
KAFKA_BOOTSTRAP_SERVERS_CONS = 'localhost:9092'
```

Each producer are assigned to the ‘counts’ topic

```
PROBLEMS 10  OUTPUT  DEBUG CONSOLE  TERMINAL  JUPYTER: VARIABLES

Preparing message: 2
Message:  {'timeuuid_id': '1663101824.6978', 'lgu_code': '1200', 'sensor_id': 'sensor_01', 'date_saved': '09/14/2022', 'tim
e_saved': '04:43:44', 'total': 9, 'car': 2, 'bus': 2, 'truck': 2, 'jeepney': 0, 'bike': 2, 'tryke': 1, 'others': 0}
Preparing message: 3
Message:  {'timeuuid_id': '1663101825.69972', 'lgu_code': '1200', 'sensor_id': 'sensor_01', 'date_saved': '09/14/2022', 'ti
me_saved': '04:43:45', 'total': 11, 'car': 3, 'bus': 2, 'truck': 1, 'jeepney': 2, 'bike': 0, 'tryke': 2, 'others': 1}
Preparing message: 4
Message:  {'timeuuid_id': '1663101826.701794', 'lgu_code': '1200', 'sensor_id': 'sensor_01', 'date_saved': '09/14/2022', 't
ime_saved': '04:43:46', 'total': 14, 'car': 4, 'bus': 2, 'truck': 2, 'jeepney': 2, 'bike': 1, 'tryke': 3, 'others': 0}
```

Producer 1

```
PROBLEMS 10  OUTPUT  DEBUG CONSOLE  TERMINAL  JUPYTER: VARIABLES

Preparing message: 2
Message:  {'timeuuid_id': '1663101855.611263', 'lgu_code': '1200', 'sensor_id': 'sensor_02', 'date_saved': '09/14/2022', 't
ime_saved': '04:44:15', 'total': 12, 'car': 4, 'bus': 0, 'truck': 0, 'jeepney': 1, 'bike': 4, 'tryke': 2, 'others': 1}
Preparing message: 3
Message:  {'timeuuid_id': '1663101856.613235', 'lgu_code': '1200', 'sensor_id': 'sensor_02', 'date_saved': '09/14/2022', 't
ime_saved': '04:44:16', 'total': 10, 'car': 0, 'bus': 1, 'truck': 1, 'jeepney': 0, 'bike': 4, 'tryke': 2, 'others': 2}
Preparing message: 4
Message:  {'timeuuid_id': '1663101857.615473', 'lgu_code': '1200', 'sensor_id': 'sensor_02', 'date_saved': '09/14/2022', 't
ime_saved': '04:44:17', 'total': 9, 'car': 4, 'bus': 2, 'truck': 0, 'jeepney': 0, 'bike': 1, 'tryke': 1, 'others': 1}
```

Producer 2

```
PROBLEMS  OUTPUT  DEBUG CONSOLE  TERMINAL  JUPYTER

thon/debugpy/adapter/../../debugpy/launcher 41921 -- /home/cara/CS172/producer3.py
Kafka Producer Application Started ...
Preparing message: 1
Message:  {'timeuuid_id': '1663085282.220497', 'lgu_code': '1200', 'sensor_id': 'sensor_03', 'date_saved': '09/14/2022', 't
ime_saved': '00:08:02', 'total': 5, 'car': 1, 'bus': 1, 'truck': 0, 'jeepney': 1, 'bike': 2, 'tryke': 0, 'others': 0}
Preparing message: 2
Message:  {'timeuuid_id': '1663085283.228356', 'lgu_code': '1200', 'sensor_id': 'sensor_03', 'date_saved': '09/14/2022', 't
ime_saved': '00:08:03', 'total': 6, 'car': 0, 'bus': 1, 'truck': 2, 'jeepney': 0, 'bike': 1, 'tryke': 0, 'others': 2}
Preparing message: 3
Message:  {'timeuuid_id': '1663085284.230445', 'lgu_code': '1200', 'sensor_id': 'sensor_03', 'date_saved': '09/14/2022', 't
ime_saved': '00:08:04', 'total': 6, 'car': 2, 'bus': 0, 'truck': 1, 'jeepney': 2, 'bike': 0, 'tryke': 1, 'others': 0}
```

Producer 3

```
PROBLEMS 10  OUTPUT  DEBUG CONSOLE  TERMINAL  JUPYTER: VARIABLES

ime_saved': '04:45:46', 'total': 10, 'car': 2, 'bus': 0, 'truck': 1, 'jeepney': 1, 'bike': 4, 'tryke': 1, 'others': 1}
Preparing message: 3
Message:  {'timeuuid_id': '1663101947.931751', 'lgu_code': '1200', 'sensor_id': 'sensor_04', 'date_saved': '09/14/2022', 't
ime_saved': '04:45:47', 'total': 11, 'car': 4, 'bus': 2, 'truck': 0, 'jeepney': 0, 'bike': 1, 'tryke': 2, 'others': 2}
Preparing message: 4
Message:  {'timeuuid_id': '1663101948.934005', 'lgu_code': '1200', 'sensor_id': 'sensor_04', 'date_saved': '09/14/2022', 't
ime_saved': '04:45:48', 'total': 5, 'car': 0, 'bus': 0, 'truck': 0, 'jeepney': 2, 'bike': 2, 'tryke': 0, 'others': 1}
Preparing message: 5
Message:  {'timeuuid_id': '1663101949.937408', 'lgu_code': '1200', 'sensor_id': 'sensor_04', 'date_saved': '09/14/2022', 't
ime_saved': '04:45:49', 'total': 9, 'car': 3, 'bus': 2, 'truck': 1, 'jeepney': 0, 'bike': 2, 'tryke': 0, 'others': 1}
```

Producer 4

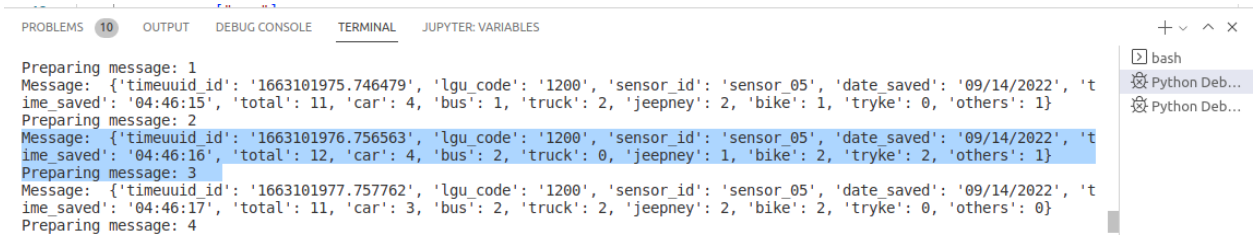

Producer 5

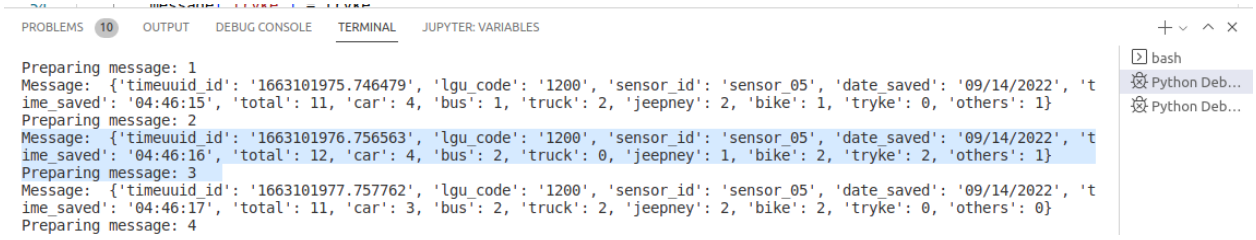

Producer 6

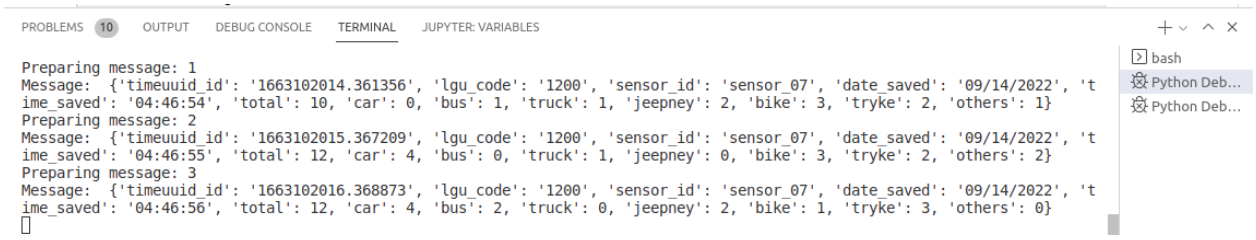

Producer 7

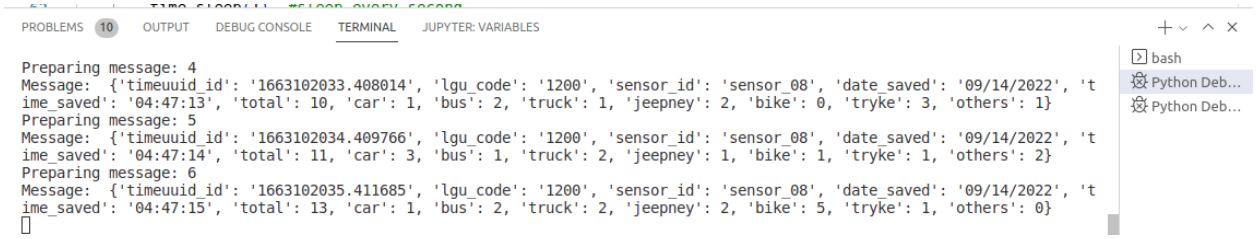

Producer 8

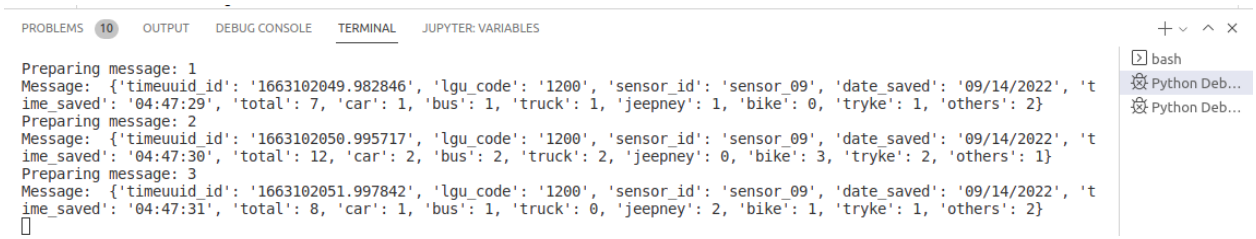

Producer 9

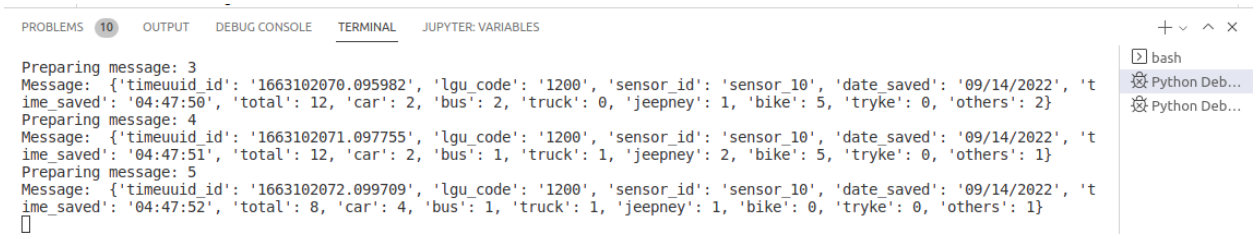

Producer 10

Part III. Data read by the consumer

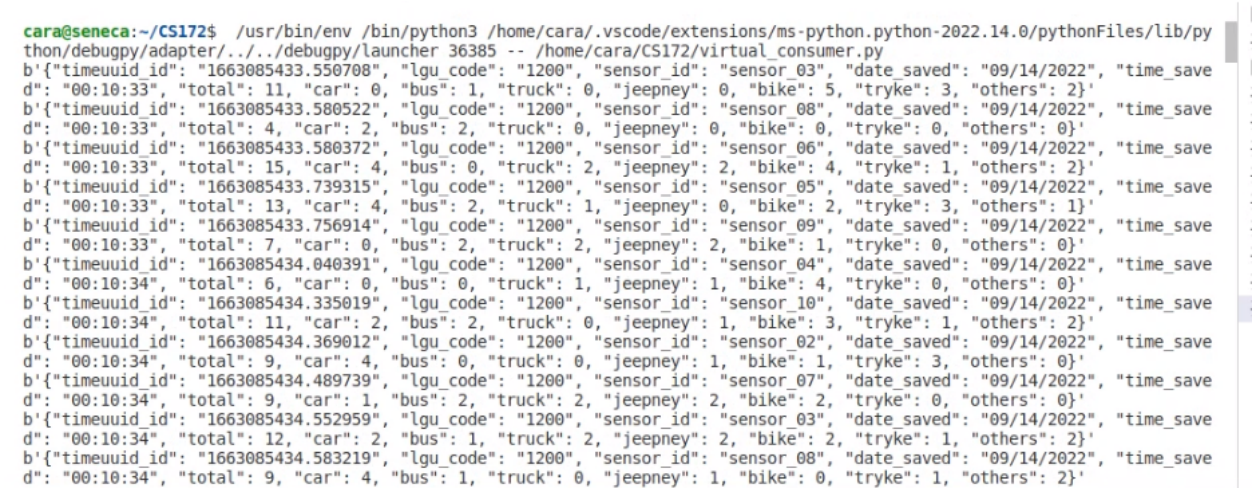

Data from different producers read by consumer

Part IV. Create a database and a table in Cassandra

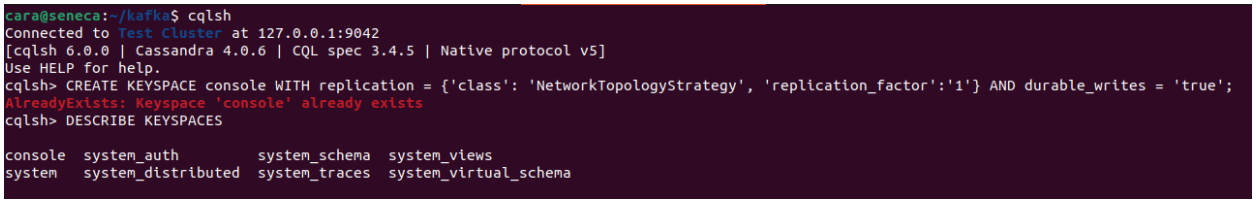

Created Keyspace ‘console’.

```
cara@seneca:~/kafka$ cqlsh
Connected to Test Cluster at 127.0.0.1:9042
[cqlsh 6.0.0 | Cassandra 4.0.6 | CQL spec 3.4.5 | Native protocol v5]
Use HELP for help.
cqlsh> USE console
   ... ;
cqlsh:console> CREATE TABLE cctv_vehicle_counts (timeuuid_id text PRIMARY KEY, lgu_code text, sensor_id text, data_saved text, time_saved text, total
int, car int, bus int, truck int, jeepney int, bike int, tryke int, others int);
AlreadyExists: Table 'console.cctv_vehicle_counts' already exists
cqlsh:console>  DESCRIBE TABLE cctv_vehicle_counts

CREATE TABLE console.cctv_vehicle_counts (
    timeuuid_id text PRIMARY KEY,
    bike int,
    bus int,
    car int,
    data_saved text,
    jeepney int,
    lgu_code text,
    others int,
    sensor_id text,
    time_saved text,
    total int,
    truck int,
    tryke int
) WITH additional_write_policy = '99p'
    AND bloom_filter_fp_chance = 0.01
    AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
    AND cdc = false
    AND comment = ''
    AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
    AND compression = {'chunk_length_in_kb': '16', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
    AND crc_check_chance = 1.0
    AND default_time_to_live = 0
```

Created a table in the ‘console’ database named ‘cctv_vehicle_counts’

```
cqlsh:console> select * from console.cctv_vehicle_counts
           ... ;

 timeuuid_id | bike | bus | car | data_saved | jeepney | lgu_code | others | sensor_id | time_saved | total | truck | tryke
-------------+------+-----+-----+------------+---------+----------+--------+-----------+------------+-------+-------+-------

(0 rows)
cqlsh:console>
cqlsh:console>
```

Check if table exists

Part V. Python program. Kafka to Cassandra

Python code. Gets access to the cassandra database then loads values into it from kafka. Consumer sink is used here, a kafka connector.

```
bootstrap_servers = ['localhost:9092']
topicname = 'counts'
consumer = KafkaConsumer(topicname, bootstrap_servers =
bootstrap_servers,auto_offset_reset='earliest',group_id="test-consumer-group")
cluster = Cluster(['127.0.0.1'], port = 9042)
session = cluster.connect('console')
```

This is the code to get access to the servers, topics, and cassandra.

py | producer9.py | Notes | virtual_consumer.py | kafka_consumer_s

Dockerfile

```
FROM python:3
WORKDIR /home/cara/CS172
COPY requirements.txt ./
RUN pip install --no-cache-dir -r requirements.txt
COPY...
CMD [ "python", "-u", "./kafka_consumer_sink.py" ]
```

Dockerfile to run the program

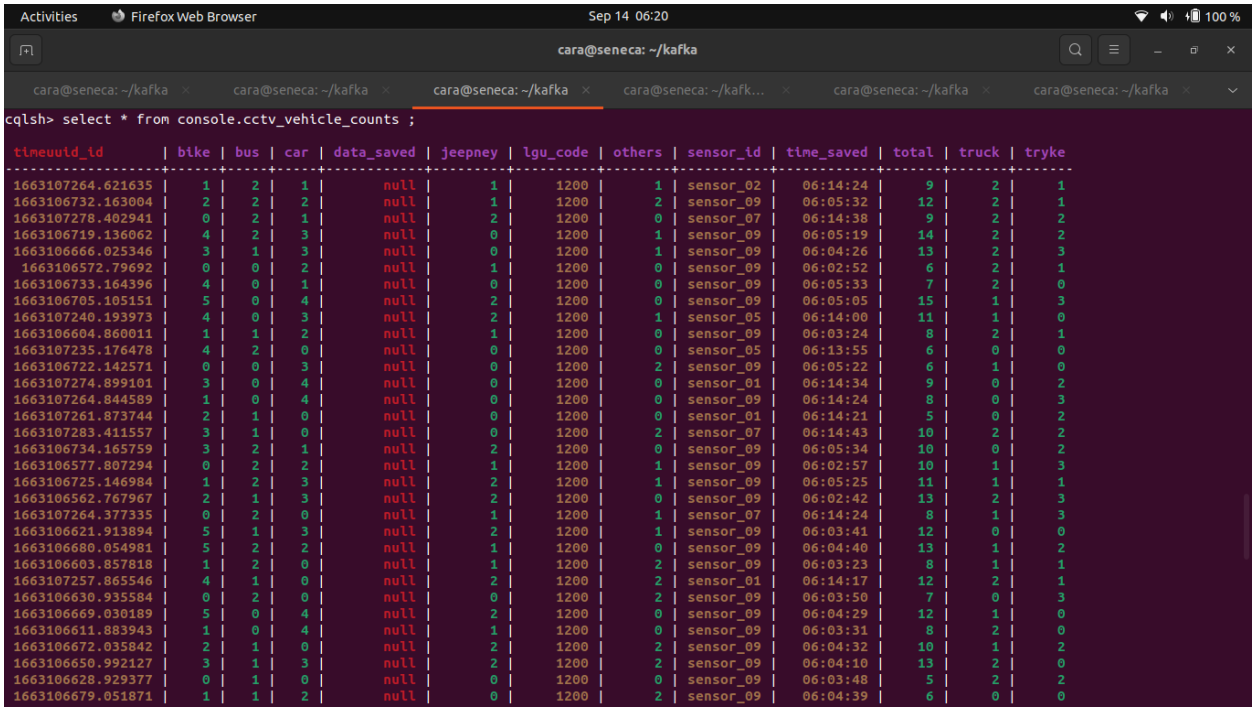

timeuuid_id	bike	bus	car	data_saved	jeepney	lgu_code	others	sensor_id	time_saved	total	truck	tryke
1663107264.621635	1	2	1	null	1	1200	1	sensor_02	06:14:24	9	2	1
1663106732.163004	2	2	2	null	1	1200	2	sensor_09	06:05:32	12	2	1
1663107278.402941	0	2	1	null	2	1200	0	sensor_07	06:14:38	9	2	2
1663106719.136062	4	2	3	null	0	1200	1	sensor_09	06:05:19	14	2	2
1663106666.025346	3	1	3	null	0	1200	1	sensor_09	06:04:26	13	2	3
1663106572.79692	0	0	2	null	1	1200	0	sensor_09	06:02:52	6	2	1
1663106733.164396	4	0	1	null	0	1200	0	sensor_09	06:05:33	7	2	0
1663106705.105151	5	0	4	null	2	1200	0	sensor_09	06:05:05	15	1	3
1663107240.193973	4	0	3	null	2	1200	1	sensor_05	06:14:00	11	1	0
1663106604.860011	1	1	2	null	1	1200	0	sensor_09	06:03:24	8	2	1
1663107235.176478	4	2	0	null	0	1200	0	sensor_05	06:13:55	6	0	0
1663106722.142571	0	0	3	null	0	1200	2	sensor_09	06:05:22	6	1	0
1663107274.899101	3	0	4	null	0	1200	0	sensor_01	06:14:34	9	0	2
1663107264.844589	1	0	4	null	0	1200	0	sensor_09	06:14:24	8	0	3
1663107261.873744	2	1	0	null	0	1200	0	sensor_01	06:14:21	5	0	2
1663107283.411557	3	1	0	null	0	1200	2	sensor_07	06:14:43	10	2	2
1663106734.165759	3	2	1	null	2	1200	0	sensor_09	06:05:34	10	0	2
1663106577.807294	0	2	2	null	1	1200	1	sensor_09	06:02:57	10	1	3
1663106725.146984	1	2	3	null	2	1200	1	sensor_09	06:05:25	11	1	1
1663106562.767967	2	1	3	null	2	1200	0	sensor_09	06:02:42	13	2	3
1663107264.377335	0	2	0	null	1	1200	1	sensor_07	06:14:24	8	1	3
1663106621.913894	5	1	3	null	2	1200	1	sensor_09	06:03:41	12	0	0
1663106680.054981	5	2	2	null	1	1200	0	sensor_09	06:04:40	13	1	2
1663106603.857818	1	2	0	null	1	1200	2	sensor_09	06:03:23	8	1	1
1663107257.865546	4	1	0	null	2	1200	2	sensor_01	06:14:17	12	2	1
1663106630.935584	0	2	0	null	0	1200	2	sensor_09	06:03:50	7	0	3
1663106669.030189	5	0	4	null	2	1200	0	sensor_09	06:04:29	12	1	0
1663106611.883943	1	0	4	null	1	1200	0	sensor_09	06:03:31	8	2	0
1663106672.035842	2	1	0	null	2	1200	2	sensor_09	06:04:32	10	1	2
1663106650.992127	3	1	3	null	2	1200	2	sensor_09	06:04:10	13	2	0
1663106628.929377	0	1	0	null	0	1200	0	sensor_09	06:03:48	5	2	2
1663106679.051871	1	1	2	null	0	1200	2	sensor_09	06:04:39	6	0	0

After running different producers the ‘kafka_consumer_sink.py’ file will be run. The data from kafka will show in the table ‘cctv_vehicle_counts’ in the ‘console’ keyspace. The image above outputs data from different producers using the *‘select *’* command.